

Deeper Networks for Image Classification

ECS795P Deep Learning and
Computer Vision
Jahnvi Sikligar
210538601

Abstract— Image Classification is a difficult task that involves categorizing images. To solve the image classification problem, deep learning models are frequently used. The goal of this paper is to investigate various Deep CNN architectures by implementing RedNet50 and VGG16 and perform image classification on MNIST and CIFAR10 datasets. For evaluation purpose both the architectures have been trained with similar epochs and batch sizes.

Keywords—Image classification, ResNet50, VGG16, MNIST, CIFAR10, CNN

I. INTRODUCTION

Deep convolutional neural networks have made numerous advances in image classification. Deep convolution networks typically incorporate low/mid/high level highlights and classifiers in an end-to-end multilayer style, and the "levels" of features can be advanced by the number of stacked layers (depth). Fortunately, current GPUs, combined with exceptionally optimised 2D convolution implementations, are powerful enough to enable the training of large CNNs, and ongoing datasets, such as ImageNet, contain enough labelled examples to train models without significant overfitting. The following are the specific contributions to this report: implementation of two deep networks, VGG and ResNet. Deep neural networks were trained using the datasets MNIST and CIFAR10. In this paper, we review the experiments performed with the two models VGG and ResNet on the MNIST and CIFAR10 datasets, where the datasets and testing results are discussed.

II. CRITICAL ANALYSIS

AlexNet was the first to use Convolutional Networks in the field of Computer Vision. Its architecture was very similar to LeNet's, but it was deeper, larger, and featured Convolutional Layers stacked on top of each other instead of pooling layers after each convolution layer. ZFNet was created by altering the architecture hyperparameters, such as lowering the stride and filter size of the first layer and increasing the size of the convolutional layers in the middle of the network.

VGGNet, the runner-up in the 2014 ILSVRC, was introduced by Karen Simonyan and Andrew Zisserman. This network emphasised the importance of network depth in achieving good performance. VGG16, their final best network, has convolution/fully-connected layers and an extremely homogeneous architecture that only performs 3x3 convolutions and 2x2 pooling throughout the network. One of the issues with the VGGNet is the large number of parameters (140M), the majority of which are in the FC layer, which requires a lot more memory and is very expensive.

GoogLeNet, which was launched by Google in 2014, was the ILSVRC 2014 winning architecture. GoogLeNet introduced the concept of Inception Module, which dramatically reduced the number of parameters in the network (4M, compared to AlexNet's 60M) despite increasing the network's depth and width.

GoogLeNet overcame VGGNet's limitations by using the Average Pooling layer instead of the Fully Connected layers at the top of the ConvNet, reducing the number of parameters without affecting network performance. So far, many improved variants of GoogLeNet have been released, the most recent being Inception-v4. Deep model training takes time, and due to a lack of data, such models are prone to overfitting.

ResNet (Residual Network) architecture was introduced in 2015 to improve training while making networks deeper. The degradation problem occurs as networks become deeper. The architecture uses special skip connections to connect output from the previous layer to the layer ahead, making the network deeper while ensuring that the error rate is not higher than the architecture's shallower versions.

III. METHOD / MODEL DESCRIPTION

A. VGG16

i. Description

VGGNet arose from the need to reduce the number of parameters in convolution layers and improve training efficiency. It is an example of a CNN that focuses on spatial exploitation. One of the main disadvantages of VGG is that it uses approximately 13 million parameters. To help reduce the number of variables, VGG employs a fixed filter of size 3x3 in the hidden layers. It is a 19-layered network that successfully demonstrated that the simultaneous placement of small sized filters such as (3x3) could produce the same effect as a large sized filter such as (5x5) or (7x7). A max-pooling layer placed after the convolutional layer aids in network tuning.

ii. Architecture

VGG is a 19-layer deep network that performs convolutions with fixed-size kernels. VGG's original architecture accepts a 224x224 RGB image as input and routes it through a network of hidden layers. The number of hidden layers is determined by the VGG variant used. For this course, VGG16, which has 16 hidden layers is being used. The network includes a fixed 3x3 filter with stride 1, five MaxPool layers, each following a batch of convolution layers with a 2x2 kernel and stride 2. These contribute to network fine tuning. These layers have been activated with ReLU, followed by Dense layers with 4096 neurons and the final output with softmax activation.

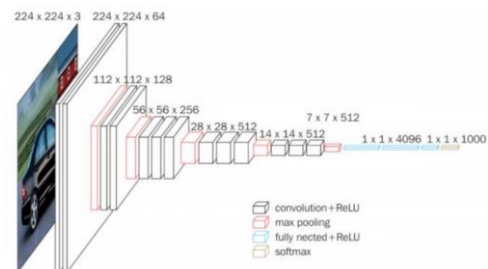


Fig A – VGG16 Architecture

B. ResNet50

i. Description

ResNet, or Residual Neural Network, is a novel architecture that features "skip connections" and heavy batch normalisation. Such skip connections are known as gated recurrent units or gated units, and they are similar to recent successful RNN elements. This method can train neural networks with 152 layers while remaining less complex than VGGNet. All ResNet configurations follow a similar configuration, with the only difference being the depth of building blocks (shown in brackets). From 18 layers (ResNet18) to 152 layers (ResNet152) (ResNet152).

ii. Architecture

The ResNet50 architecture is based on the above model, but there is one significant difference. Due to concerns about the time required to train the layers, the building block was modified into a bottleneck design in this case. Instead of the previous two layers, this time a three-layer stack was used. As a result, each of the ResNet34's 2-layer blocks was replaced with a 3-layer bottleneck block, resulting in the ResNet 50 architecture. This model is much more accurate than the 34-layer ResNet model. ResNet's 50-layer performance is 3.8 billion FLOPS.

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112	7×7, 64, stride 2				
3×3 max pool, stride 2						
conv2.x	56×56	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3.x	28×28	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4.x	14×14	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5.x	7×7	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
1×1		average pool, 1000-d fc, softmax				
FLOPs		1.8×10 ⁹	3.6×10 ⁹	3.8×10 ⁹	7.6×10 ⁹	11.3×10 ⁹

Fig B – ResNet50 Architecture

C. Datasets

i. MNIST

The MNIST (Modified National Institute of Standards and Technology) database of handwritten digits 0 to 9 contains 60,000 training examples and 10,000 test examples. Each image is 28 x 28 pixels in size. There are a total of ten classes that correspond to the digits [0-9]. All of the images are one-channel grayscale.

ii. CIFAR-10

CIFAR10 is made up of 60000 coloured images of size 32x32 (50,000 in the training set and 10,000 in the test set) (3-channels). There are ten classes, each with 60000 images. Airplane, automobile, bird, cat, deer, dog, frog, horse, ship, and truck are the class names.

D. Data Augmentation

Image rotation is a popular augmentation technique that allows the model to become insensitive to object orientation. By passing an integer value in the rotation range argument to the ImageDataGenerator class, you can randomly rotate images through any degree between 0 and 360.

E. Parameter Tuning

i. Optimizers

VGGNet and ResNet are trained using various optimizers such as Stochastic Gradient Descent, momentum addition, and the ADAM optimizer. ADAM is commonly regarded as an improvement over RMSprop for improving model performance.

ii. Epochs and Batch Size

For all models, the batch size and epochs are currently fixed at 128 and 50, respectively. To train at the maximum batch size capacity for each model, batch size is reduced. Increasing the batch size improves model accuracy at the expense of computation cost.

IV. EXPERIMENTS

The models in this paper have been trained on both MNIST and CIFAR10 datasets. The results of them have been discussed below with same configurations for all the models. Configurations:- Epochs: 50, Batch size: 128.

A. MNIST

i. Training a. VGG 16

```
# save model
model_name = model_to_json(
    open('architecture_VGG16.json', 'w').write(model_json)
    model.save_weights('weights_VGG16.h5', overwrite=True)

# list all data in history
print(history.history.keys())
print(history.history)

WARNING:tensorflow:From /tensorflow-1.15.1/python3.7/tensorflow/tensorflow/tensorflow.py:422: The name tf.global_variables is deprecated.
Epoch 1/50
48000/48000 [=====] - 32s 660us/step - loss: 2.3801 - accuracy: 0.1160 - val_loss: 2.2960 - val_accuracy: 0.1860
Epoch 2/50
48000/48000 [=====] - 27s 550us/step - loss: 2.2300 - accuracy: 0.2180 - val_loss: 1.9392 - val_accuracy: 0.2860
Epoch 3/50
48000/48000 [=====] - 27s 550us/step - loss: 0.6963 - accuracy: 0.7724 - val_loss: 0.1342 - val_accuracy: 0.9405
Epoch 4/50
48000/48000 [=====] - 27s 550us/step - loss: 0.1762 - accuracy: 0.9474 - val_loss: 0.1276 - val_accuracy: 0.9620
Epoch 5/50
48000/48000 [=====] - 27s 550us/step - loss: 0.1375 - accuracy: 0.9650 - val_loss: 0.1389 - val_accuracy: 0.9684
Epoch 6/50
48000/48000 [=====] - 27s 550us/step - loss: 0.4055 - accuracy: 0.9773 - val_loss: 0.0620 - val_accuracy: 0.9807
Epoch 7/50
48000/48000 [=====] - 26s 550us/step - loss: 0.0589 - accuracy: 0.9820 - val_loss: 0.0768 - val_accuracy: 0.9778
Epoch 8/50
48000/48000 [=====] - 26s 550us/step - loss: 0.0465 - accuracy: 0.9858 - val_loss: 0.0710 - val_accuracy: 0.9801
Epoch 9/50
48000/48000 [=====] - 26s 550us/step - loss: 0.0376 - accuracy: 0.9880 - val_loss: 0.0545 - val_accuracy: 0.9835
Epoch 10/50
48000/48000 [=====] - 27s 550us/step - loss: 0.0317 - accuracy: 0.9903 - val_loss: 0.0543 - val_accuracy: 0.9852
Test score: 0.065236791632525
Test accuracy: 0.9901806139763
dict_keys(['val_loss', 'val_accuracy', 'loss', 'accuracy'])
{'val_loss': 2.2959943055315, 'val_accuracy': 0.18600000000000002, 'loss': 2.3800999999999998, 'accuracy': 0.11600000000000001}
```

b. ResNet10

```
Epoch 1/50
48000/48000 [=====] - 32s 660us/step - loss: 2.3801 - accuracy: 0.1160 - val_loss: 2.2960 - val_accuracy: 0.1860
Epoch 2/50
48000/48000 [=====] - 27s 550us/step - loss: 2.2300 - accuracy: 0.2180 - val_loss: 1.9392 - val_accuracy: 0.2860
Epoch 3/50
48000/48000 [=====] - 27s 550us/step - loss: 0.6963 - accuracy: 0.7724 - val_loss: 0.1342 - val_accuracy: 0.9405
Epoch 4/50
48000/48000 [=====] - 27s 550us/step - loss: 0.1762 - accuracy: 0.9474 - val_loss: 0.1276 - val_accuracy: 0.9620
Epoch 5/50
48000/48000 [=====] - 27s 550us/step - loss: 0.1375 - accuracy: 0.9650 - val_loss: 0.1389 - val_accuracy: 0.9684
Epoch 6/50
48000/48000 [=====] - 27s 550us/step - loss: 0.4055 - accuracy: 0.9773 - val_loss: 0.0620 - val_accuracy: 0.9807
Epoch 7/50
48000/48000 [=====] - 26s 550us/step - loss: 0.0589 - accuracy: 0.9820 - val_loss: 0.0768 - val_accuracy: 0.9778
Epoch 8/50
48000/48000 [=====] - 26s 550us/step - loss: 0.0465 - accuracy: 0.9858 - val_loss: 0.0710 - val_accuracy: 0.9801
Epoch 9/50
48000/48000 [=====] - 26s 550us/step - loss: 0.0376 - accuracy: 0.9880 - val_loss: 0.0545 - val_accuracy: 0.9835
Epoch 10/50
48000/48000 [=====] - 27s 550us/step - loss: 0.0317 - accuracy: 0.9903 - val_loss: 0.0543 - val_accuracy: 0.9852
Test score: 0.065236791632525
Test accuracy: 0.9901806139763
dict_keys(['val_loss', 'val_accuracy', 'loss', 'accuracy'])
{'val_loss': 2.2959943055315, 'val_accuracy': 0.18600000000000002, 'loss': 2.3800999999999998, 'accuracy': 0.11600000000000001}
```

- ii. Testing
 - a. VGG 16

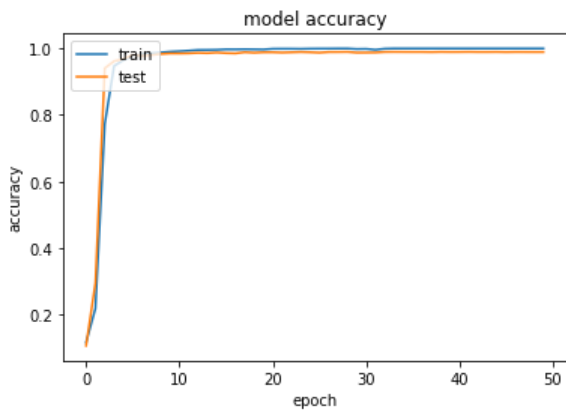


Fig. C – Graph of Testing accuracy vs Epoch

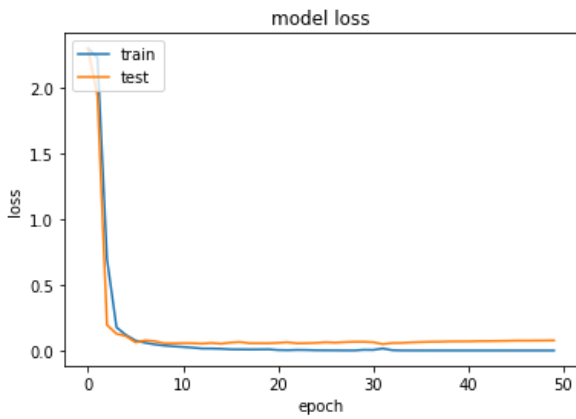


Fig. D – Graph of Testing loss vs Epoch

b. ResNet10

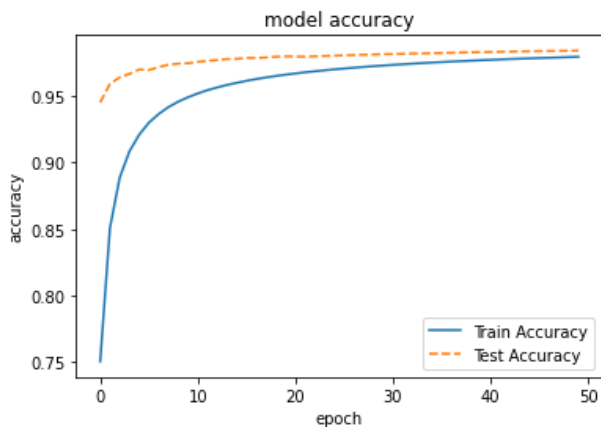


Fig. D – Graph of Testing accuracy vs Epoch

and test accuracy of 98.43% on Epoch 50. The **Table I** summarizes the ‘*Training accuracy*’ and ‘*Testing accuracy*’ of both VGG16 and ResNet50 models respectively.

TABLE I.

MNIST Dataset		
<i>Models</i>	<i>Training</i>	<i>Testing</i>
VGG16	98.9300%	99.0100 %
ResNet50	97.9701 %	98.4304 %

B. CIFAR-10

- i. Training
 - a. VGG 16

```

Train on 40000 samples, validate on 10000 samples
Epoch 1/50
40000/40000 [=====] - 279.71ms/step - loss: 2.3026 - accuracy: 0.1082 - val_loss: 2.3026 - val_accuracy: 0.1106
Epoch 2/50
40000/40000 [=====] - 279.64ms/step - loss: 2.3025 - accuracy: 0.1087 - val_loss: 2.3025 - val_accuracy: 0.1088
Epoch 3/50
40000/40000 [=====] - 279.61ms/step - loss: 2.3025 - accuracy: 0.1013 - val_loss: 2.3025 - val_accuracy: 0.0952
Epoch 4/50
40000/40000 [=====] - 275.613ms/step - loss: 2.3024 - accuracy: 0.1024 - val_loss: 2.3024 - val_accuracy: 0.0952
Epoch 5/50
40000/40000 [=====] - 276.68ms/step - loss: 2.3023 - accuracy: 0.1024 - val_loss: 2.3023 - val_accuracy: 0.0952
Epoch 6/50
40000/40000 [=====] - 267.639ms/step - loss: 2.3022 - accuracy: 0.1054 - val_loss: 2.3022 - val_accuracy: 0.0952
Epoch 7/50
40000/40000 [=====] - 269.618ms/step - loss: 2.3020 - accuracy: 0.1021 - val_loss: 2.3019 - val_accuracy: 0.0952
Epoch 8/50
40000/40000 [=====] - 269.618ms/step - loss: 2.3017 - accuracy: 0.1082 - val_loss: 2.3016 - val_accuracy: 0.0952
Epoch 9/50
40000/40000 [=====] - 269.68ms/step - loss: 2.3012 - accuracy: 0.1214 - val_loss: 2.3009 - val_accuracy: 0.0952
Epoch 10/50
40000/40000 [=====] - 269.64ms/step - loss: 2.3003 - accuracy: 0.1205 - val_loss: 2.2996 - val_accuracy: 0.1411
Epoch 11/50
40000/40000 [=====] - 269.618ms/step - loss: 2.2989 - accuracy: 0.1579 - val_loss: 2.2969 - val_accuracy: 0.1789
Epoch 12/50
40000/40000 [=====] - 269.618ms/step - loss: 2.2984 - accuracy: 0.1747 - val_loss: 2.2907 - val_accuracy: 0.1810
Epoch 13/50
40000/40000 [=====] - 269.68ms/step - loss: 2.2881 - accuracy: 0.1811 - val_loss: 2.2745 - val_accuracy: 0.1808
Testing: 4.712s

Epoch 36/50
40000/40000 [=====] - 269.64ms/step - loss: 1.4728 - accuracy: 0.4488 - val_loss: 1.4857 - val_accuracy: 0.4456
Epoch 37/50
40000/40000 [=====] - 269.64ms/step - loss: 1.4428 - accuracy: 0.4632 - val_loss: 1.4604 - val_accuracy: 0.4531
Epoch 38/50
40000/40000 [=====] - 269.64ms/step - loss: 1.4047 - accuracy: 0.4813 - val_loss: 1.3830 - val_accuracy: 0.4853
Epoch 39/50
40000/40000 [=====] - 269.64ms/step - loss: 1.3640 - accuracy: 0.4978 - val_loss: 1.4118 - val_accuracy: 0.4714
Epoch 40/50
40000/40000 [=====] - 269.64ms/step - loss: 1.3319 - accuracy: 0.5113 - val_loss: 1.4064 - val_accuracy: 0.4714
Epoch 41/50
40000/40000 [=====] - 269.64ms/step - loss: 1.3092 - accuracy: 0.5192 - val_loss: 1.2868 - val_accuracy: 0.5316
Epoch 42/50
40000/40000 [=====] - 269.64ms/step - loss: 1.2706 - accuracy: 0.5336 - val_loss: 1.2825 - val_accuracy: 0.5268
Epoch 43/50
40000/40000 [=====] - 269.64ms/step - loss: 1.2348 - accuracy: 0.5499 - val_loss: 1.2151 - val_accuracy: 0.5558
Epoch 44/50
40000/40000 [=====] - 269.64ms/step - loss: 1.2041 - accuracy: 0.5611 - val_loss: 1.4104 - val_accuracy: 0.4875
Epoch 45/50
40000/40000 [=====] - 269.64ms/step - loss: 1.1763 - accuracy: 0.5677 - val_loss: 1.1295 - val_accuracy: 0.5807
Epoch 46/50
40000/40000 [=====] - 269.64ms/step - loss: 1.1316 - accuracy: 0.5888 - val_loss: 1.1156 - val_accuracy: 0.5934
Epoch 47/50
40000/40000 [=====] - 269.64ms/step - loss: 1.0895 - accuracy: 0.5988 - val_loss: 1.0798 - val_accuracy: 0.6180
Epoch 48/50
40000/40000 [=====] - 269.64ms/step - loss: 1.0654 - accuracy: 0.6113 - val_loss: 1.0728 - val_accuracy: 0.6160
Epoch 49/50
40000/40000 [=====] - 269.64ms/step - loss: 1.0654 - accuracy: 0.6113 - val_loss: 1.0728 - val_accuracy: 0.6160
10000/10000 [=====] - 26.15ms/step

Test score: 1.080062273212476
Test accuracy: 0.606700031471252
dict key=[ 'val_loss', 'val_accuracy', 'loss', 'accuracy' ]
[ 'val_loss': [ 2.30251545318623, 2.302508962789554, 2.302471553095508, 2.30239098233427, 2.302305262713624, 2.302205166134648, 2.302105166134648, 2.302005166134648, 2.301905166134648, 2.301805166134648, 2.301705166134648, 2.301605166134648, 2.301505166134648, 2.301405166134648, 2.301305166134648, 2.301205166134648, 2.301105166134648, 2.301005166134648, 2.300905166134648, 2.300805166134648, 2.300705166134648, 2.300605166134648, 2.300505166134648, 2.300405166134648, 2.300305166134648, 2.300205166134648, 2.300105166134648, 2.300005166134648, 2.299905166134648, 2.299805166134648, 2.299705166134648, 2.299605166134648, 2.299505166134648, 2.299405166134648, 2.299305166134648, 2.299205166134648, 2.299105166134648, 2.299005166134648, 2.298905166134648, 2.298805166134648, 2.298705166134648, 2.298605166134648, 2.298505166134648, 2.298405166134648, 2.298305166134648, 2.298205166134648, 2.298105166134648, 2.298005166134648, 2.297905166134648, 2.297805166134648, 2.297705166134648, 2.297605166134648, 2.297505166134648, 2.297405166134648, 2.297305166134648, 2.297205166134648, 2.297105166134648, 2.297005166134648, 2.296905166134648, 2.296805166134648, 2.296705166134648, 2.296605166134648, 2.296505166134648, 2.296405166134648, 2.296305166134648, 2.296205166134648, 2.296105166134648, 2.296005166134648, 2.295905166134648, 2.295805166134648, 2.295705166134648, 2.295605166134648, 2.295505166134648, 2.295405166134648, 2.295305166134648, 2.295205166134648, 2.295105166134648, 2.295005166134648, 2.294905166134648, 2.294805166134648, 2.294705166134648, 2.294605166134648, 2.294505166134648, 2.294405166134648, 2.294305166134648, 2.294205166134648, 2.294105166134648, 2.294005166134648, 2.293905166134648, 2.293805166134648, 2.293705166134648, 2.293605166134648, 2.293505166134648, 2.293405166134648, 2.293305166134648, 2.293205166134648, 2.293105166134648, 2.293005166134648, 2.292905166134648, 2.29280516613
```

ii. Testing
a. VGG 16

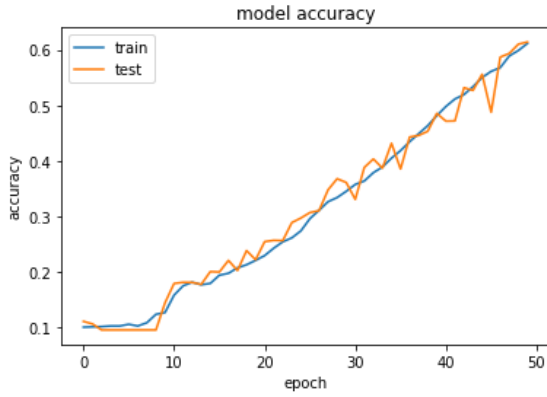


Fig. E – Graph of Testing accuracy vs Epoch

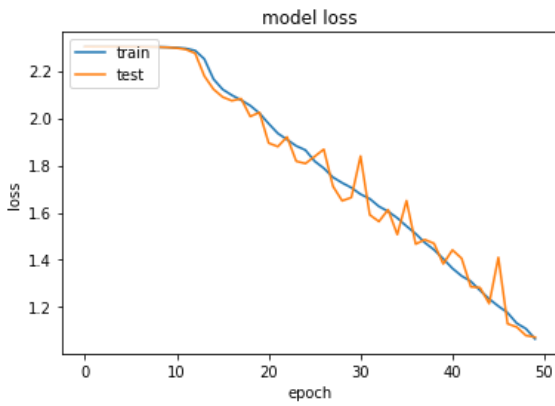


Fig. F – Graph of Testing accuracy vs Epoch

b. ResNet10

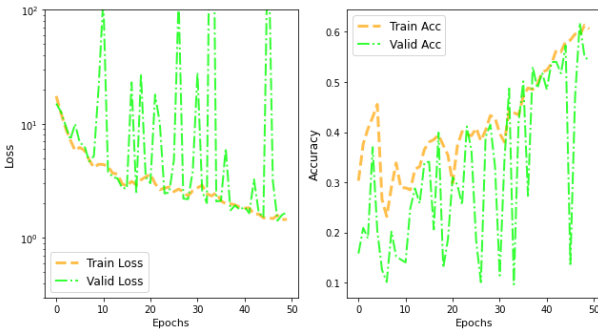


Fig. G – Graph of Testing accuracy vs Epoch (right)
Fig. H – Graph of Testing loss vs Epoch (left)
Note: 'Validation' here refers to 'Testing'

Classification Report:

	precision	recall	f1-score	support
airplane	0.22	0.84	0.35	1000
automobile	0.69	0.60	0.64	1000
bird	0.73	0.31	0.44	1000
cat	0.67	0.02	0.04	1000
deer	0.81	0.14	0.23	1000
dog	0.00	0.00	0.00	1000
frog	0.49	0.87	0.63	1000
horse	0.55	0.76	0.64	1000
ship	0.85	0.43	0.57	1000
truck	0.74	0.74	0.74	1000
accuracy			0.47	10000
macro avg	0.57	0.47	0.43	10000
weighted avg	0.57	0.47	0.43	10000

Fig. I – Further evaluation

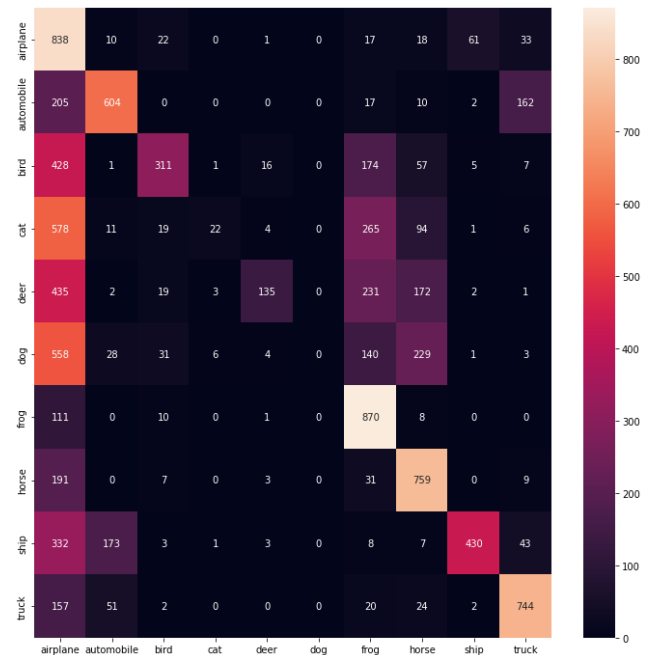


Fig. J – Heatmap for Confusion Matrix

According to the environment settings and execution output, the VGG16 model achieved training accuracy of 61.40% and test accuracy of 60.67% on Epoch 50. The **Table II** summarizes the 'Training accuracy' and 'Testing accuracy' of both VGG16 and ResNet50 models respectively.

TABLE II.

CIFAR-10 Dataset		
Models	Training	Testing
VGG16	61.4000 %	60.6700 %
ResNet50	54.1900 %	53.9200 %

V. CONCLUSION

We investigated the performance of VGG16 and ResNet50 networks on different datasets while using the same parameter tuning. These networks, despite their size, make

full use of the CNN architecture of DNN. Dense networks like these have an impact on the model's performance and run time. The results show that both the VGG-16 and ResNet50 deep convolutional neural networks can achieve high accuracy results on the MNIST dataset using purely supervised learning. To keep the experiment simple, both models were tested with MNIST datasets to see how well they performed. In comparison to the VGG-16 model, the ResNet50 model has a faster training rate and better early epoch results.

FURTHER EVALUATION

It will be interesting to compare the performance of the vgg-16 and resnet50 models on a difficult dataset such as ImageNet in the future. Increasing the layer count to upgrade to vgg19 and ResNet152 models and observing performance with various types of datasets. Alternately, you can evaluate model performance by adding or removing a layer.

REFERENCES

- [1] Krizhevsky, A., Sutskever, I. and Hinton, G.E., 2012. Imagenet classification with deep convolutional neural networks. In *Advances in neural information processing systems* (pp. 1097-1105).
- [2] Rawat, W. and Wang, Z., 2017. Deep convolutional neural networks for image classification: A comprehensive review. *Neural computation*, 29(9), pp.2352-2449.
- [3] Zeiler, M.D. and Fergus, R., 2014, September. Visualizing and understanding convolutional networks. In *European conference on computer vision* (pp. 818-833). Springer, Cham.
- [4] Simonyan, K. and Zisserman, A., 2014. Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*.
- [5] He, K., Zhang, X., Ren, S. and Sun, J., 2016. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 770-778).